

dataxlab-internship-task5

November 7, 2025

```
[1]: # O. Imports & settings
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
```

```
[2]: # plotting settings
%matplotlib inline
sns.set(style="whitegrid")
pd.set_option('display.max_columns', 200)
```

```
[3]: import pandas as pd
file_path = r"D:\intern\Superstore.csv"

df = pd.read_csv(file_path, encoding='latin1')
print(df.head())
```

```
Row ID          Order ID  Order Date   Ship Date      Ship Mode Customer ID \
0     1 CA-2016-152156  11/8/2016  11/11/2016 Second Class CG-12520
1     2 CA-2016-152156  11/8/2016  11/11/2016 Second Class CG-12520
2     3 CA-2016-138688  6/12/2016  6/16/2016 Second Class DV-13045
3     4 US-2015-108966 10/11/2015 10/18/2015 Standard Class SO-20335
4     5 US-2015-108966 10/11/2015 10/18/2015 Standard Class SO-20335
```

```
Customer Name    Segment      Country           City      State \
0 Claire Gute    Consumer United States Henderson Kentucky
1 Claire Gute    Consumer United States Henderson Kentucky
2 Darrin Van Huff Corporate United States Los Angeles California
3 Sean O'Donnell Consumer United States Fort Lauderdale Florida
4 Sean O'Donnell Consumer United States Fort Lauderdale Florida
```

```
Postal Code Region       Product ID      Category Sub-Category \
0        42420   South FUR-B0-10001798 Furniture Bookcases
1        42420   South FUR-CH-10000454 Furniture Chairs
2        90036   West OFF-LA-10000240 Office Supplies Labels
```

```

3      33311  South  FUR-TA-10000577      Furniture      Tables
4      33311  South  OFF-ST-10000760  Office Supplies     Storage

                                         Product Name      Sales  Quantity \
0          Bush Somerset Collection Bookcase  261.9600           2
1  Hon Deluxe Fabric Upholstered Stacking Chairs,...  731.9400           3
2  Self-Adhesive Address Labels for Typewriters b...   14.6200           2
3      Bretford CR4500 Series Slim Rectangular Table  957.5775           5
4      Eldon Fold 'N Roll Cart System    22.3680           2

  Discount      Profit
0      0.00  41.9136
1      0.00  219.5820
2      0.00   6.8714
3      0.45 -383.0310
4      0.20   2.5164

```

```
[4]: # 2. Quick overview
df.info()
df.head()
df.describe(include='all').T
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9994 entries, 0 to 9993
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Row ID             9994 non-null   int64  
 1   Order ID           9994 non-null   object  
 2   Order Date         9994 non-null   object  
 3   Ship Date          9994 non-null   object  
 4   Ship Mode          9994 non-null   object  
 5   Customer ID        9994 non-null   object  
 6   Customer Name       9994 non-null   object  
 7   Segment             9994 non-null   object  
 8   Country             9994 non-null   object  
 9   City                9994 non-null   object  
 10  State               9994 non-null   object  
 11  Postal Code         9994 non-null   int64  
 12  Region              9994 non-null   object  
 13  Product ID          9994 non-null   object  
 14  Category             9994 non-null   object  
 15  Sub-Category         9994 non-null   object  
 16  Product Name         9994 non-null   object  
 17  Sales                9994 non-null   float64 
 18  Quantity             9994 non-null   int64  
 19  Discount             9994 non-null   float64 
 20  Profit               9994 non-null   float64 

```

dtypes: float64(3), int64(3), object(15)
memory usage: 1.6+ MB

	count	unique	top	freq	mean	\
Row ID	9994.0	NaN	NaN	NaN	4997.5	
Order ID	9994	5009	CA-2017-100111	14	NaN	
Order Date	9994	1237	9/5/2016	38	NaN	
Ship Date	9994	1334	12/16/2015	35	NaN	
Ship Mode	9994	4	Standard Class	5968	NaN	
Customer ID	9994	793	WB-21850	37	NaN	
Customer Name	9994	793	William Brown	37	NaN	
Segment	9994	3	Consumer	5191	NaN	
Country	9994	1	United States	9994	NaN	
City	9994	531	New York City	915	NaN	
State	9994	49	California	2001	NaN	
Postal Code	9994.0	NaN	NaN	NaN	55190.379428	
Region	9994	4	West	3203	NaN	
Product ID	9994	1862	OFF-PA-10001970	19	NaN	
Category	9994	3	Office Supplies	6026	NaN	
Sub-Category	9994	17	Binders	1523	NaN	
Product Name	9994	1850	Staple envelope	48	NaN	
Sales	9994.0	NaN	NaN	NaN	229.858001	
Quantity	9994.0	NaN	NaN	NaN	3.789574	
Discount	9994.0	NaN	NaN	NaN	0.156203	
Profit	9994.0	NaN	NaN	NaN	28.656896	
	std	min	25%	50%	75%	max
Row ID	2885.163629	1.0	2499.25	4997.5	7495.75	9994.0
Order ID	NaN	NaN	NaN	NaN	NaN	NaN
Order Date	NaN	NaN	NaN	NaN	NaN	NaN
Ship Date	NaN	NaN	NaN	NaN	NaN	NaN
Ship Mode	NaN	NaN	NaN	NaN	NaN	NaN
Customer ID	NaN	NaN	NaN	NaN	NaN	NaN
Customer Name	NaN	NaN	NaN	NaN	NaN	NaN
Segment	NaN	NaN	NaN	NaN	NaN	NaN
Country	NaN	NaN	NaN	NaN	NaN	NaN
City	NaN	NaN	NaN	NaN	NaN	NaN
State	NaN	NaN	NaN	NaN	NaN	NaN
Postal Code	32063.69335	1040.0	23223.0	56430.5	90008.0	99301.0
Region	NaN	NaN	NaN	NaN	NaN	NaN
Product ID	NaN	NaN	NaN	NaN	NaN	NaN
Category	NaN	NaN	NaN	NaN	NaN	NaN
Sub-Category	NaN	NaN	NaN	NaN	NaN	NaN
Product Name	NaN	NaN	NaN	NaN	NaN	NaN
Sales	623.245101	0.444	17.28	54.49	209.94	22638.48
Quantity	2.22511	1.0	2.0	3.0	5.0	14.0
Discount	0.206452	0.0	0.0	0.2	0.2	0.8

```
Profit           234.260108 -6599.978  1.72875   8.6665   29.364   8399.976
```

```
[5]: # 3. Missing values & duplicates
missing = df.isnull().sum().sort_values(ascending=False)
missing[missing>0]
print("Duplicates:", df.duplicated().sum())
```

```
Duplicates: 0
```

```
[6]: # 4. Value counts for categorical fields (pick a few important ones)
for col in df.select_dtypes(include='object').columns[:8]:
    print("----", col, "----")
    print(df[col].value_counts(dropna=False).head(10))
    print()
```

```
---- Order ID ----
```

```
Order ID
```

```
CA-2017-100111      14
CA-2017-157987      12
CA-2016-165330      11
US-2016-108504      11
US-2015-126977      10
CA-2016-105732      10
CA-2015-131338      10
CA-2015-158421       9
CA-2014-106439       9
US-2015-163433       9
Name: count, dtype: int64
```

```
---- Order Date ----
```

```
Order Date
```

```
9/5/2016        38
9/2/2017        36
11/10/2016      35
12/1/2017        34
12/2/2017        34
12/9/2017        33
11/12/2017      30
12/8/2017        30
9/9/2017         29
9/4/2017         28
Name: count, dtype: int64
```

```
---- Ship Date ----
```

```
Ship Date
```

```
12/16/2015      35
9/26/2017        34
11/21/2017      32
```

```
12/6/2017      32
9/6/2017       30
12/12/2017     30
9/15/2017      30
9/13/2014      27
9/8/2017       27
9/26/2015      26
Name: count, dtype: int64
```

---- Ship Mode ----

```
Ship Mode
Standard Class    5968
Second Class      1945
First Class        1538
Same Day           543
Name: count, dtype: int64
```

---- Customer ID ----

```
Customer ID
WB-21850         37
MA-17560         34
JL-15835         34
PP-18955         34
CK-12205         32
JD-15895         32
EH-13765         32
SV-20365         32
ZC-21910         31
EP-13915         31
Name: count, dtype: int64
```

---- Customer Name ----

```
Customer Name
William Brown      37
Matt Abelman       34
John Lee            34
Paul Prost          34
Chloris Kastensmidt 32
Jonathan Doherty    32
Edward Hooks        32
Seth Vernon          32
Zuschuss Carroll     31
Emily Phan            31
Name: count, dtype: int64
```

---- Segment ----

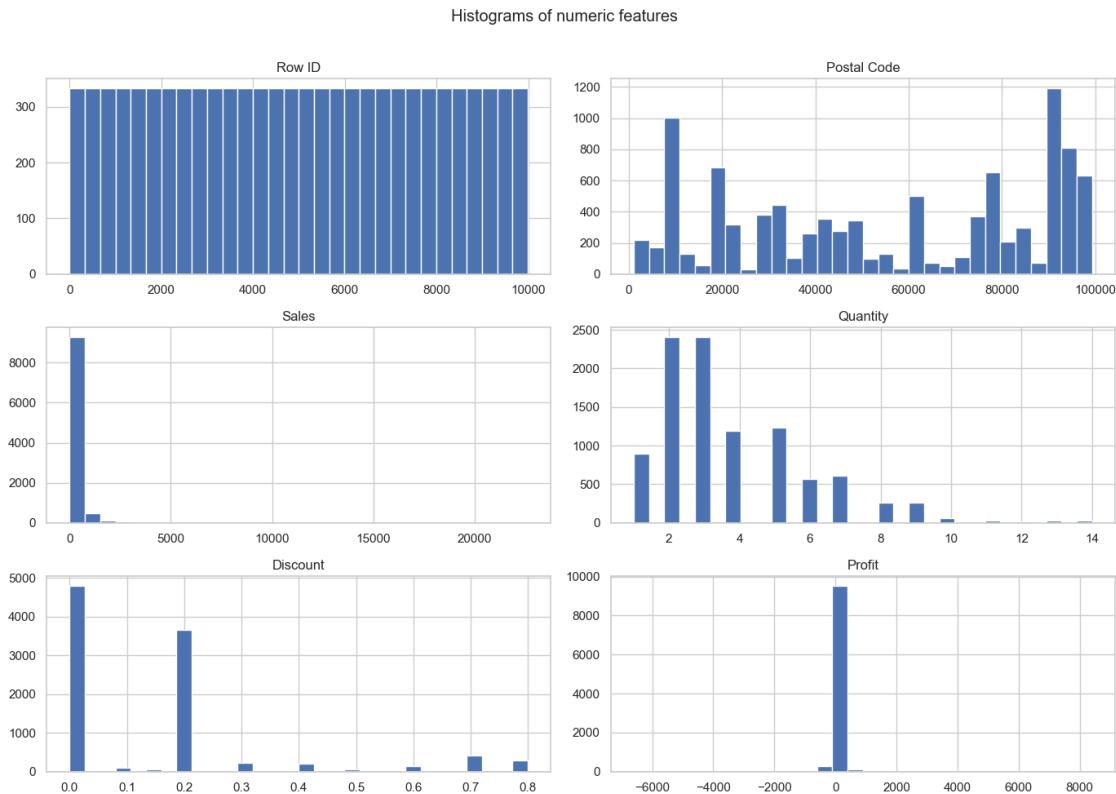
```
Segment
Consumer           5191
```

```
Corporate      3020
Home Office    1783
Name: count, dtype: int64
```

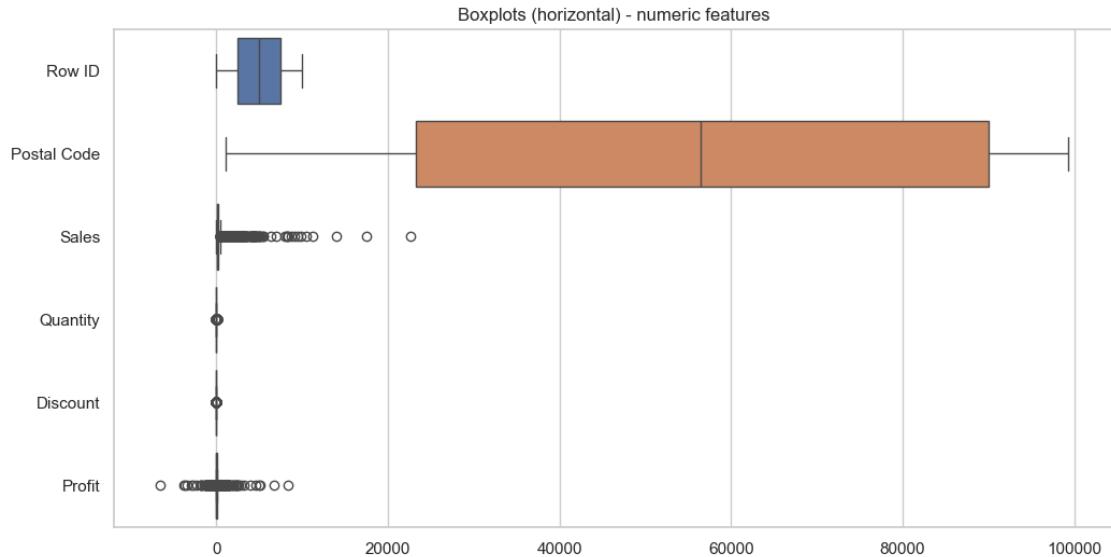
---- Country ----

```
Country
United States   9994
Name: count, dtype: int64
```

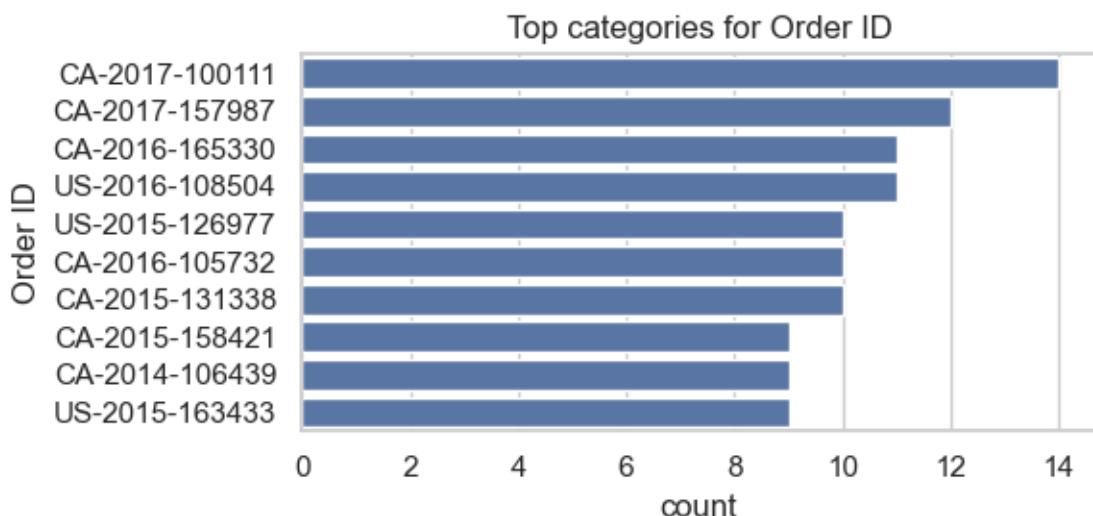
```
[7]: # 5. Univariate analysis - numeric features: histograms + boxplots
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
# histograms
df[num_cols].hist(figsize=(14, 10), bins=30)
plt.suptitle("Histograms of numeric features")
plt.tight_layout(rect=[0, 0, 1, 0.97])
```



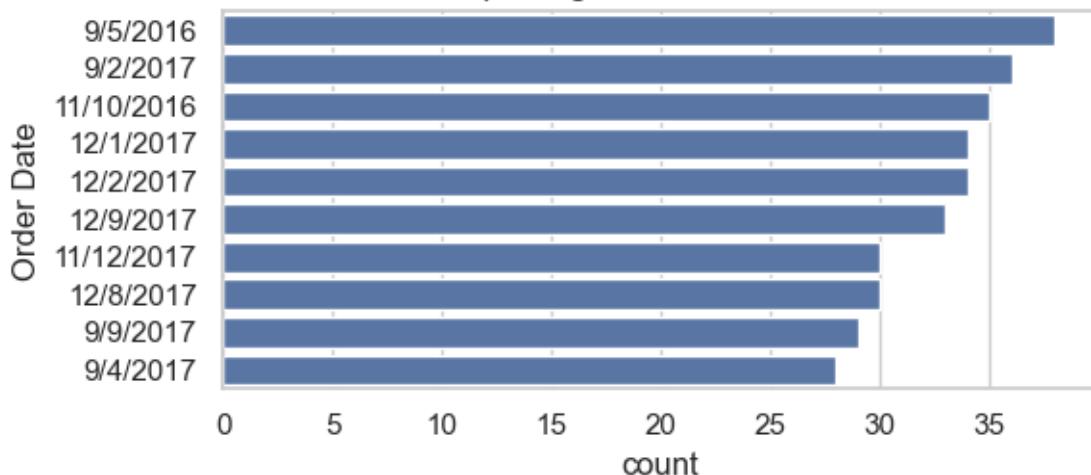
```
[8]: # boxplots (one per numeric col)
plt.figure(figsize=(12,6))
sns.boxplot(data=df[num_cols], orient='h')
plt.title("Boxplots (horizontal) - numeric features")
plt.show()
```



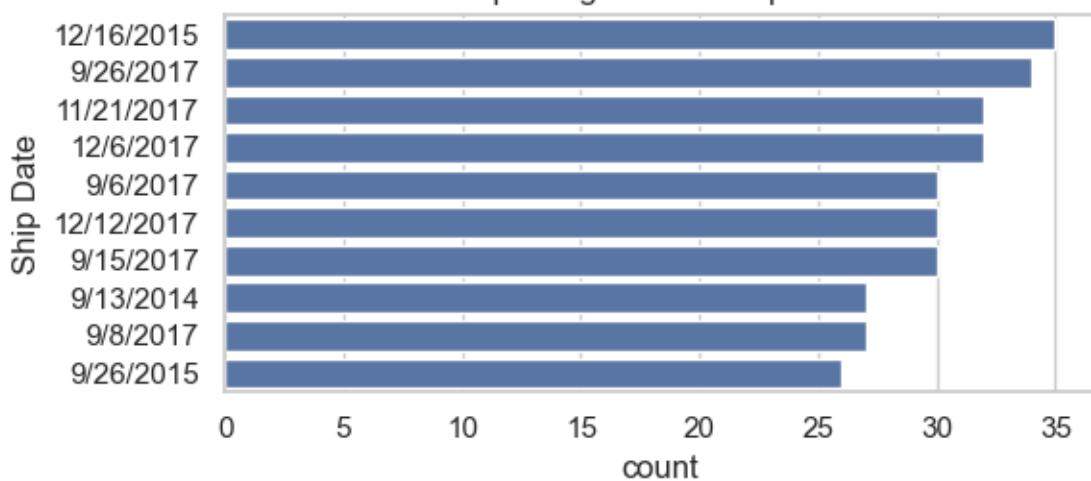
```
[9]: # 6. Categorical distributions - bar plots for top categories
cat_cols = df.select_dtypes(include='object').columns.tolist()
for c in cat_cols[:6]:
    plt.figure(figsize=(6,3))
    sns.countplot(y=c, data=df, order=df[c].value_counts().index[:10])
    plt.title(f"Top categories for {c}")
    plt.tight_layout()
    plt.show()
```

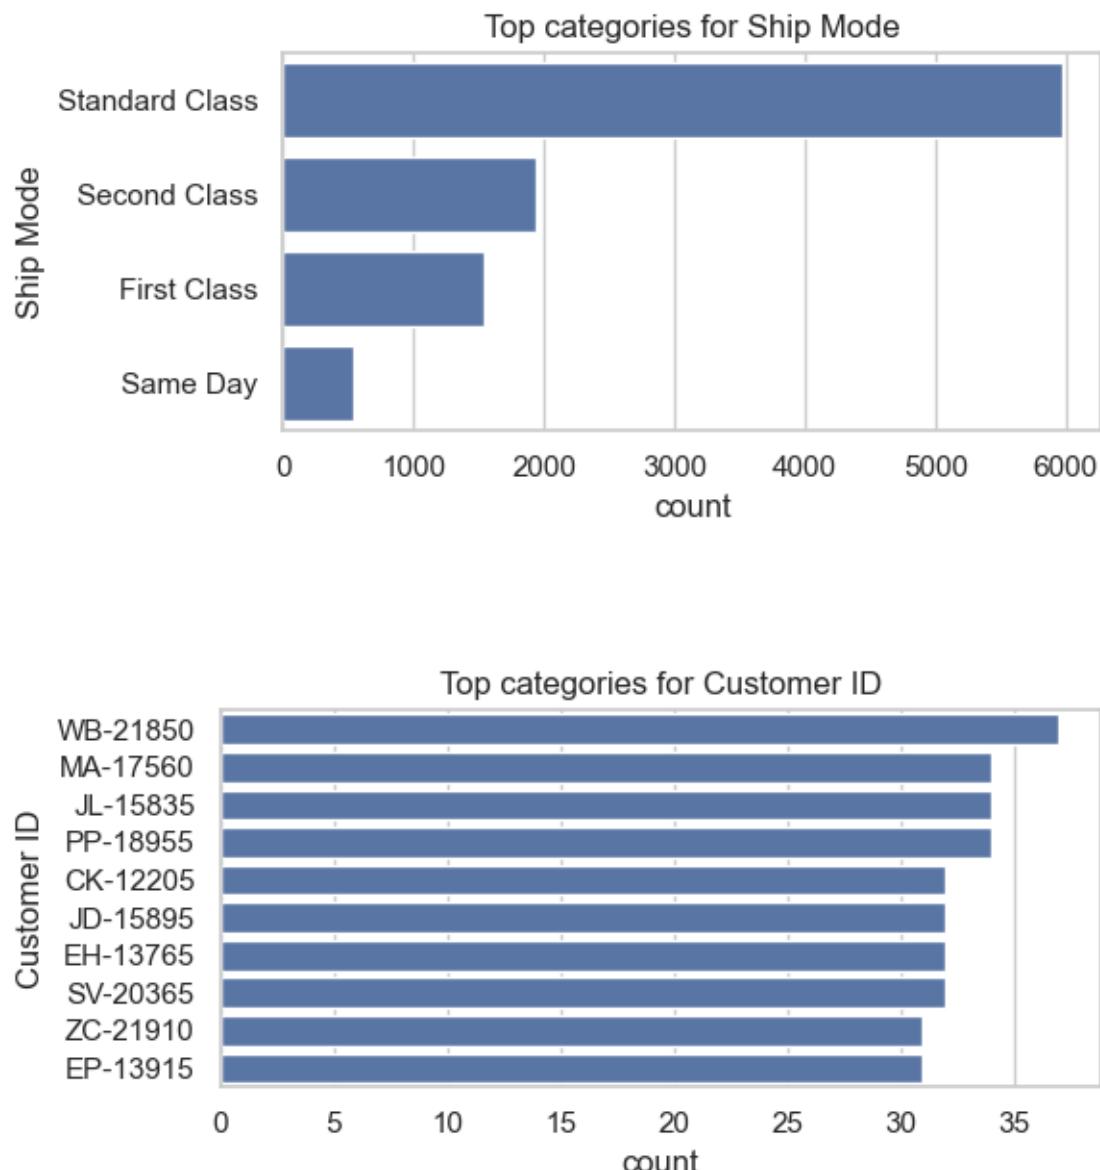


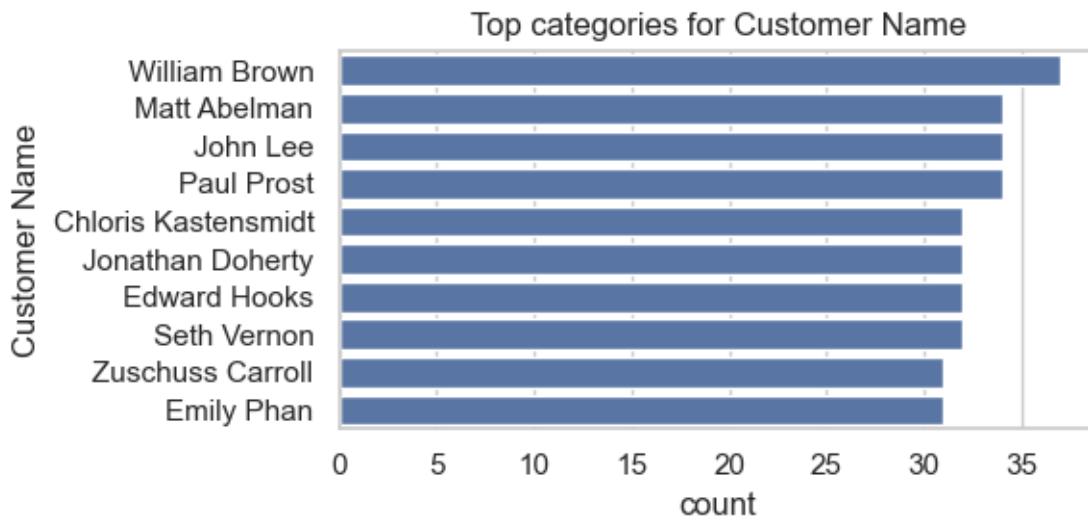
Top categories for Order Date



Top categories for Ship Date

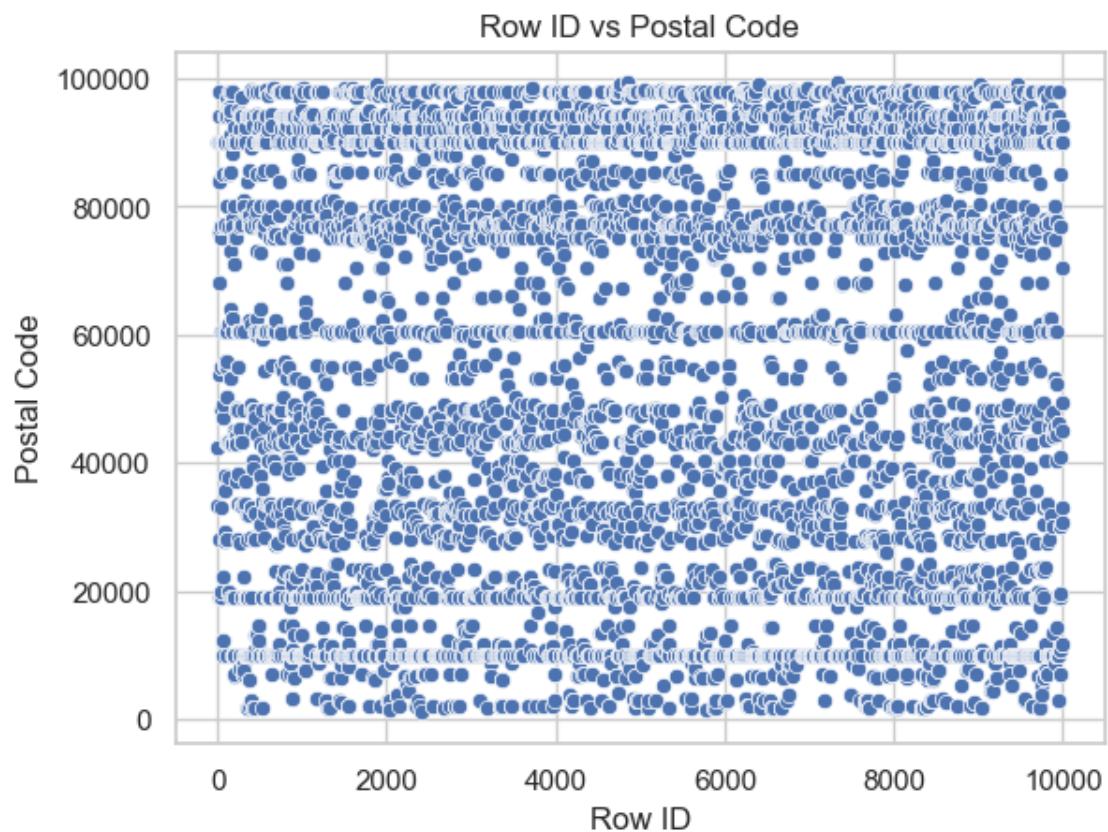


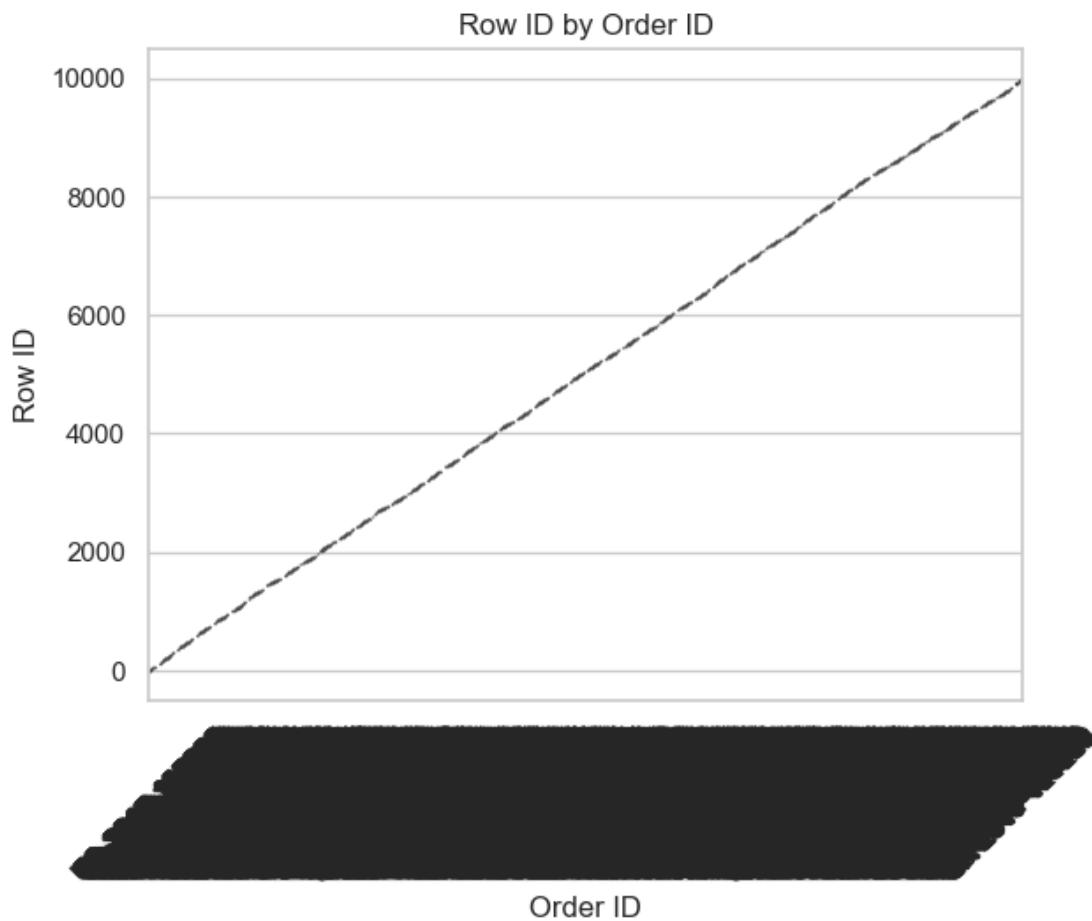




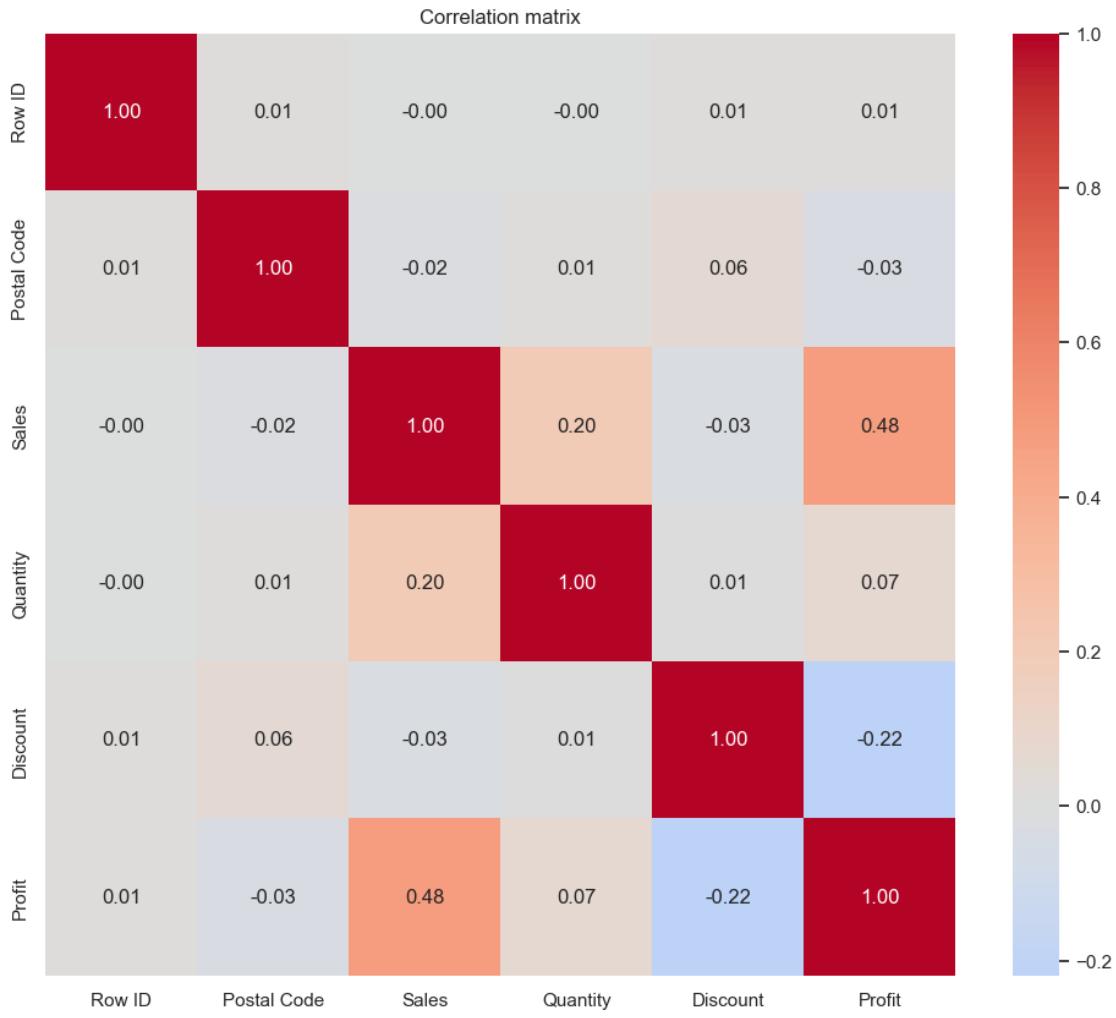
```
[10]: # 7. Bivariate analysis - scatter / boxplots
#numeric vs numeric scatter
if len(num_cols) >= 2:
    sns.scatterplot(x=num_cols[0], y=num_cols[1], data=df)
    plt.title(f'{num_cols[0]} vs {num_cols[1]}' )
    plt.show()

# numeric vs categorical: use boxplot
if cat_cols:
    sns.boxplot(x=cat_cols[0], y=num_cols[0], data=df)
    plt.title(f'{num_cols[0]} by {cat_cols[0]}' )
    plt.xticks(rotation=45)
    plt.show()
```

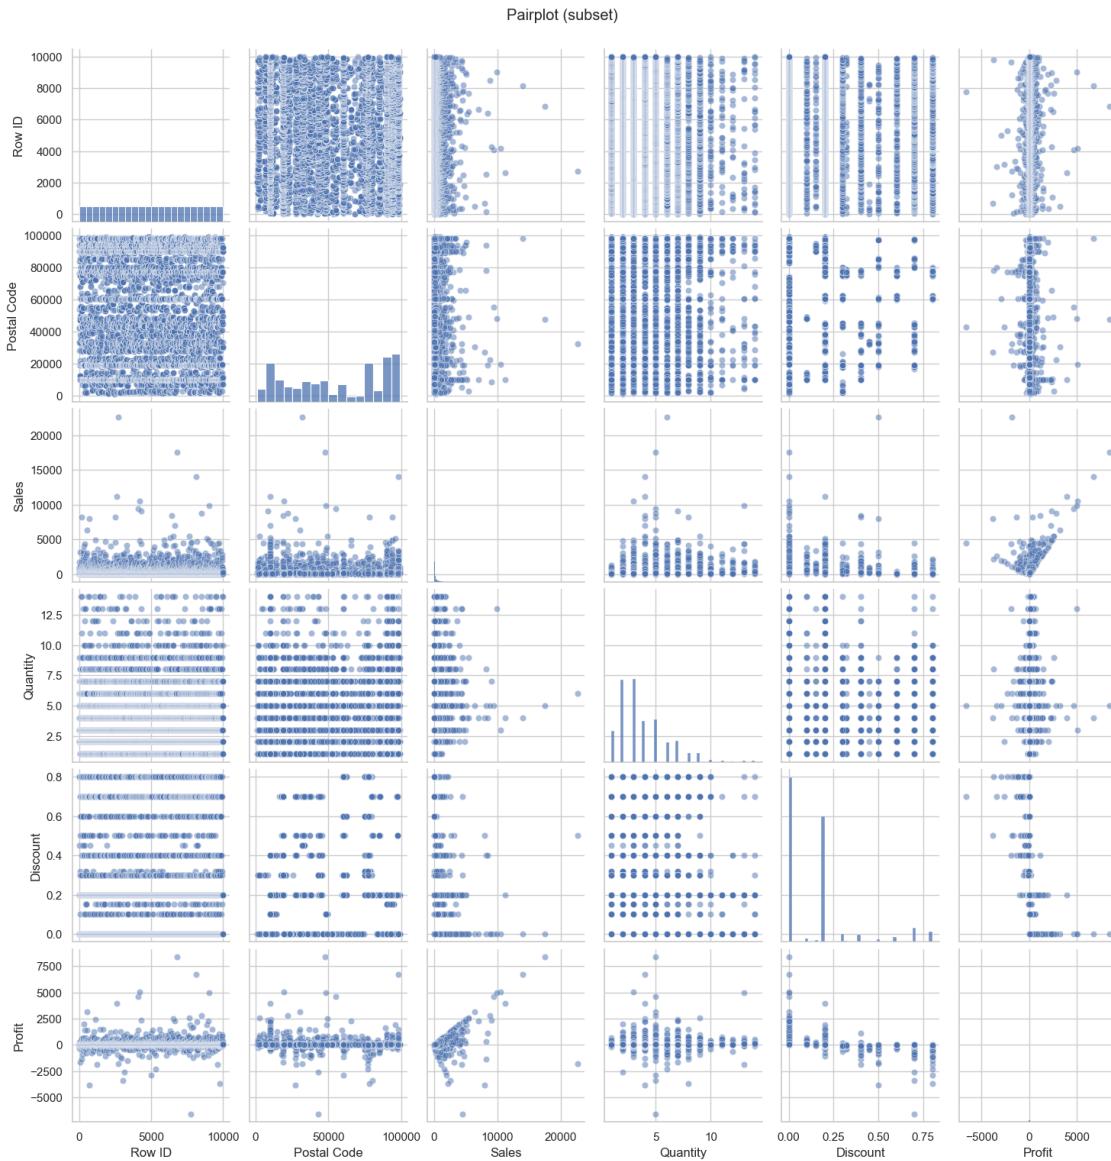




```
[11]: # 8. Correlation matrix + heatmap
corr = df[num_cols].corr()
plt.figure(figsize=(12,10))
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm', center=0)
plt.title("Correlation matrix")
plt.show()
```



```
[12]: # 9. Pairplot (careful with many columns; select subset)
subset = num_cols[:6]
sns.pairplot(df[subset].dropna(), diag_kind='hist', plot_kws={'alpha':0.5})
plt.suptitle("Pairplot (subset)", y=1.02)
plt.show()
```



```
[13]: # 10. Detect multicollinearity using VIF (for numeric features)
X = df[num_cols].dropna() # drop rows with NA in numeric columns for VIF calc
X_const = sm.add_constant(X)
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X_const.values, i+1) for i in range(len(X.columns))]
vif_data.sort_values("VIF", ascending=False)
```

	feature	VIF
5	Profit	1.375543
2	Sales	1.358932

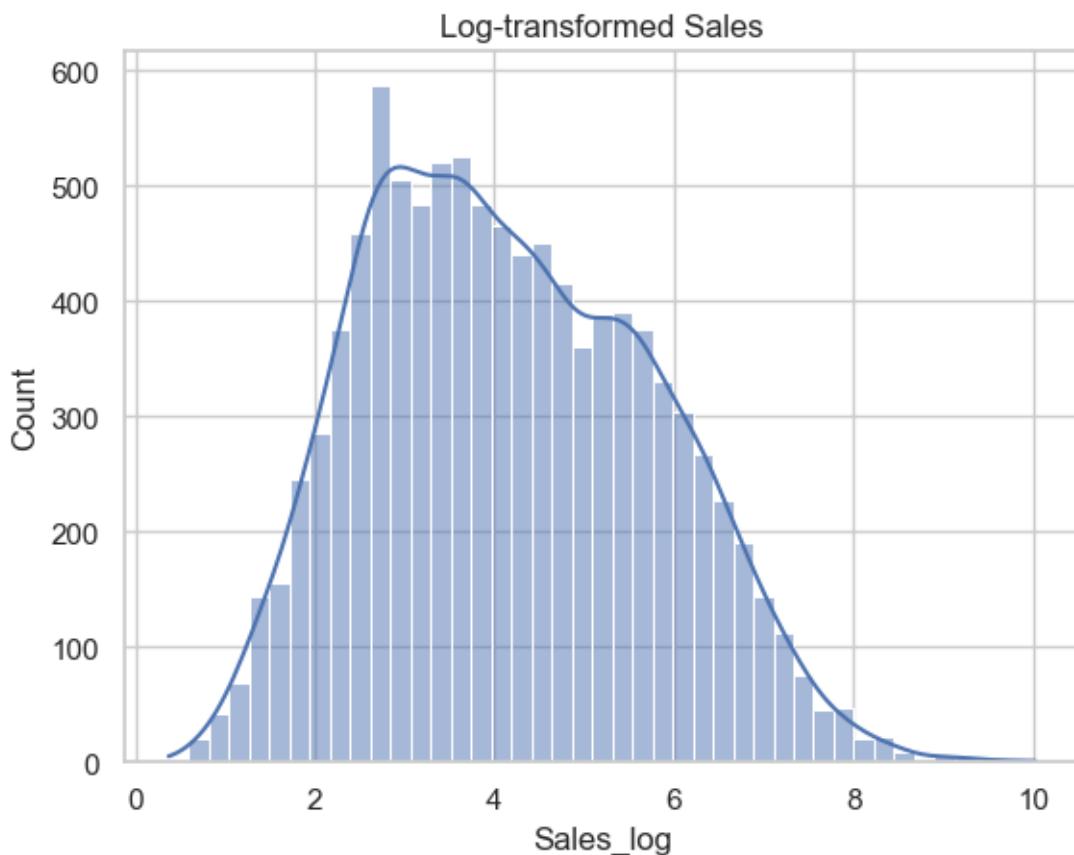
```

4      Discount  1.062648
3      Quantity   1.043633
1  Postal Code  1.004354
0      Row ID    1.000625

```

```
[14]: # 11. Skewness - show skew and example transform
skews = df[num_cols].skew().sort_values(key=lambda x: x.abs(), ascending=False)
skews.head(20)

#transform for a skewed feature
skewed_col = skews.index[0]
df[skewed_col + "_log"] = np.log1p(df[skewed_col])
sns.histplot(df[skewed_col + "_log"].dropna(), kde=True)
plt.title(f"Log-transformed {skewed_col}")
plt.show()
```



```
[15]: # 12. Outlier detection (IQR method) - count outliers per numeric col
outlier_counts = {}
for c in num_cols:
    Q1 = df[c].quantile(0.25)
    Q3 = df[c].quantile(0.75)
```

```

IQR = Q3 - Q1
lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR
outlier_counts[c] = ((df[c] < lower) | (df[c] > upper)).sum()
pd.Series(outlier_counts).sort_values(ascending=False).head(10)

```

[15]:

Profit	1881
Sales	1167
Discount	856
Quantity	170
Postal Code	0
Row ID	0
dtype:	int64

[16]:

```

# 13. Statistical test examples (e.g., t-test / ANOVA) depending on problem
#t-test for numeric across binary categorical variable
if len(cat_cols)>0 and df[cat_cols[0]].nunique() == 2:
    grp = df[df[cat_cols[0]] == df[cat_cols[0]].unique()[0]][num_cols[0]].
    ↵dropna()
    grp2 = df[df[cat_cols[0]] == df[cat_cols[0]].unique()[1]][num_cols[0]].
    ↵dropna()
    tstat, pval = stats.ttest_ind(grp, grp2, equal_var=False)
    print("t-stat:", tstat, "p-value:", pval)

```

[17]:

```

# 14. Save cleaned sample or notebook outputs if needed
df.to_csv("cleaned_dataset_sample.csv", index=False)

```

[18]:

```

# 15. Export notebook to PDF (run from terminal or Jupyter):
!jupyter nbconvert --to pdf "Dataxlab_internship_task5.ipynb"

```

This application is used to convert notebook files (*.ipynb)
to various other formats.

WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.

```

Options
=====
The options below are convenience aliases to configurable class-options,
as listed in the "Equivalent to" description-line of the aliases.
To see all configurable class-options for some <cmd>, use:
  <cmd> --help-all

--debug
  set log level to logging.DEBUG (maximize logging output)
  Equivalent to: [--Application.log_level=10]
--show-config
  Show the application's configuration (human-readable format)

```

```

Equivalent to: [--Application.show_config=True]
--show-config-json
    Show the application's configuration (json format)
Equivalent to: [--Application.show_config_json=True]
--generate-config
    generate default config file
Equivalent to: [--JupyterApp.generate_config=True]
-y
    Answer yes to any questions instead of prompting.
    Equivalent to: [--JupyterApp.answer_yes=True]
--execute
    Execute the notebook prior to export.
    Equivalent to: [--ExecutePreprocessor.enabled=True]
--allow-errors
    Continue notebook execution even if one of the cells throws an error and
    include the error message in the cell output (the default behaviour is to abort
    conversion). This flag is only relevant if '--execute' was specified, too.
    Equivalent to: [--ExecutePreprocessor.allow_errors=True]
--stdin
    read a single notebook file from stdin. Write the resulting notebook with
    default basename 'notebook.*'
    Equivalent to: [--NbConvertApp.from_stdin=True]
--stdout
    Write notebook output to stdout instead of files.
    Equivalent to: [--NbConvertApp.writer_class=StdoutWriter]
--inplace
    Run nbconvert in place, overwriting the existing notebook (only
        relevant when converting to notebook format)
    Equivalent to: [--NbConvertApp.use_output_suffix=False]
--NbConvertApp.export_format=notebook --FilesWriter.build_directory=
--clear-output
    Clear output of current file and save in place,
        overwriting the existing notebook.
    Equivalent to: [--NbConvertApp.use_output_suffix=False]
--NbConvertApp.export_format=notebook --FilesWriter.build_directory=
--ClearOutputPreprocessor.enabled=True]
--coalesce-streams
    Coalesce consecutive stdout and stderr outputs into one stream (within each
    cell).
    Equivalent to: [--NbConvertApp.use_output_suffix=False]
--NbConvertApp.export_format=notebook --FilesWriter.build_directory=
--CoalesceStreamsPreprocessor.enabled=True]
--no-prompt
    Exclude input and output prompts from converted document.
    Equivalent to: [--TemplateExporter.exclude_input_prompt=True]
--TemplateExporter.exclude_output_prompt=True]
--no-input
    Exclude input cells and output prompts from converted document.

```

This mode is ideal for generating code-free reports.

Equivalent to: [--TemplateExporter.exclude_output_prompt=True
 --TemplateExporter.exclude_input=True
 --TemplateExporter.exclude_input_prompt=True]
 --allow-chromium-download

Whether to allow downloading chromium if no suitable version is found on the system.

Equivalent to: [--WebPDFExporter.allow_chromium_download=True]
 --disable-chromium-sandbox

Disable chromium security sandbox when converting to PDF..

Equivalent to: [--WebPDFExporter.disable_sandbox=True]
 --show-input

Shows code input. This flag is only useful for dejavu users.

Equivalent to: [--TemplateExporter.exclude_input=False]
 --embed-images

Embed the images as base64 dataurls in the output. This flag is only useful for the HTML/WebPDF/Slides exports.

Equivalent to: [--HTMLExporter.embed_images=True]
 --sanitize-html

Whether the HTML in Markdown cells and cell outputs should be sanitized..

Equivalent to: [--HTMLExporter.sanitize_html=True]
 --log-level=<Enum>

Set the log level by value or name.

Choices: any of [0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR', 'CRITICAL']

Default: 30

Equivalent to: [--Application.log_level]
 --config=<Unicode>

Full path of a config file.

Default: ''

Equivalent to: [--JupyterApp.config_file]
 --to=<Unicode>

The export format to be used, either one of the built-in formats

['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook',
 'pdf', 'python', 'qtpdf', 'qtpng', 'rst', 'script', 'slides', 'webpdf']

or a dotted object name that represents the import path for an
 ``Exporter`` class

Default: ''

Equivalent to: [--NbConvertApp.export_format]
 --template=<Unicode>

Name of the template to use

Default: ''

Equivalent to: [--TemplateExporter.template_name]
 --template-file=<Unicode>

Name of the template file to use

Default: None

Equivalent to: [--TemplateExporter.template_file]
 --theme=<Unicode>

Template specific theme(e.g. the name of a JupyterLab CSS theme distributed as prebuilt extension for the lab template)
 Default: 'light'
 Equivalent to: [--HTMLExporter.theme]

--sanitize_html=<Bool>
 Whether the HTML in Markdown cells and cell outputs should be sanitized.This should be set to True by nbviewer or similar tools.
 Default: False
 Equivalent to: [--HTMLExporter.sanitize_html]

--writer=<DottedObjectName>
 Writer class used to write the results of the conversion
 Default: 'FileWriter'
 Equivalent to: [--NbConvertApp.writer_class]

--post=<DottedOrNone>
 PostProcessor class used to write the results of the conversion
 Default: ''
 Equivalent to: [--NbConvertApp.postprocessor_class]

--output=<Unicode>
 Overwrite base name use for output files.
 Supports pattern replacements '{notebook_name}'.
 Default: '{notebook_name}'
 Equivalent to: [--NbConvertApp.output_base]

--output-dir=<Unicode>
 Directory to write output(s) to. Defaults to output to the directory of each notebook.
 To recover previous default behaviour (outputting to the current working directory) use . as the flag value.
 Default: ''
 Equivalent to: [--FileWriter.build_directory]

--reveal-prefix=<Unicode>
 The URL prefix for reveal.js (version 3.x).
 This defaults to the reveal CDN, but can be any url pointing to a copy of reveal.js.
 For speaker notes to work, this must be a relative path to a local copy of reveal.js: e.g., "reveal.js".
 If a relative path is given, it must be a subdirectory of the current directory (from which the server is run).
 See the usage documentation (<https://nbconvert.readthedocs.io/en/latest/usage.html#reveal-js-html-slideshow>) for more details.
 Default: ''
 Equivalent to: [--SlidesExporter.reveal_url_prefix]

```
--nbformat=<Enum>
The nbformat version to write.
    Use this to downgrade notebooks.
Choices: any of [1, 2, 3, 4]
Default: 4
Equivalent to: [--NotebookExporter.nbformat_version]
```

Examples

The simplest way to use nbconvert is

```
> jupyter nbconvert mynotebook.ipynb --to html
```

Options include ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'qtpdf', 'qtpng', 'rst', 'script', 'slides', 'webpdf'] .

```
> jupyter nbconvert --to latex mynotebook.ipynb
```

Both HTML and LaTeX support multiple output templates. LaTeX includes

'base', 'article' and 'report'. HTML includes 'basic', 'lab' and 'classic'. You can specify the flavor of the format used.

```
> jupyter nbconvert --to html --template lab mynotebook.ipynb
```

You can also pipe the output to stdout, rather than a file

```
> jupyter nbconvert mynotebook.ipynb --stdout
```

PDF is generated via latex

```
> jupyter nbconvert mynotebook.ipynb --to pdf
```

You can get (and serve) a Reveal.js-powered slideshow

```
> jupyter nbconvert myslides.ipynb --to slides --post serve
```

Multiple notebooks can be given at the command line in a couple of different ways:

```
> jupyter nbconvert notebook*.ipynb
```

```
> jupyter nbconvert notebook1.ipynb notebook2.ipynb
```

or you can specify the notebooks list in a config file, containing::

```
c.NbConvertApp.notebooks = ["my_notebook.ipynb"]
```

```
> jupyter nbconvert --config mycfg.py
```

To see all available configurables, use `--help-all`.

```
[NbConvertApp] WARNING | pattern 'Dataxlab_internship_task5.ipynb' matched no  
files
```